

DonorsChoose

DonorsChoose.org receives hundreds of thousands of project proposals each year for classroom projects in need of funding. Right now, a large number of volunteers is needed to manually screen each submission before it's approved to be posted on the DonorsChoose.org website.

Next year, DonorsChoose.org expects to receive close to 500,000 project proposals. As a result, there are three main problems they need to solve:

- How to scale current manual processes and resources to screen 500,000 projects so that they can be posted as quickly and as efficiently as possible
- How to increase the consistency of project vetting across different volunteers to improve the experience for teachers
- How to focus volunteer time on the applications that need the most assistance

The goal of the competition is to predict whether or not a DonorsChoose.org project proposal submitted by a teacher will be approved, using the text of project descriptions as well as additional metadata about the project, teacher, and school. DonorsChoose.org can then use this information to identify projects most likely to need further review before approval.

About the DonorsChoose Data Set

The `train.csv` data set provided by DonorsChoose contains the following features:

Feature	
<code>project_id</code>	A unique identifier for the proposed project.
<code>project_title</code>	Title of the project. Art, Will, etc.
<code>project_grade_category</code>	Grade level of students for which the project is targeted.

Feature	Description
<code>project_subject_categories</code>	One or more (comma-separated) subject categories from the following enumeration: <ul style="list-style-type: none"> • Arts • English Language Arts • Foreign Languages • Health • History • Mathematics • Science • Social Studies • Technology • Visual Arts • Literacy & Language
<code>school_state</code>	State where school is located (Two-letter U.S. state abbreviations)
<code>project_subject_subcategories</code>	One or more (comma-separated) subject subcategories from the following enumeration: <ul style="list-style-type: none"> • Literature & Writing • Mathematics • Science • Social Studies • Visual Arts
<code>project_resource_summary</code>	An explanation of the resources needed for the project. <ul style="list-style-type: none"> • My students need hands on literacy materials.
<code>project_essay_1</code>	First essay
<code>project_essay_2</code>	Second essay
<code>project_essay_3</code>	Third essay
<code>project_essay_4</code>	Fourth essay
<code>project_submitted_datetime</code>	Datetime when project application was submitted. Example: 2015-01-01 12:00:00
<code>teacher_id</code>	A unique identifier for the teacher of the project. Example: bdf8baa8fedef6bfb3
<code>teacher_prefix</code>	Teacher's title. One of the following: <ul style="list-style-type: none"> • Mr. • Mrs. • Ms. • Dr. • Prof.
<code>teacher_number_of_previously_posted_projects</code>	Number of project applications previously submitted by the teacher

* See the section **Notes on the Essay Data** for more details about these features.

Additionally, the `resources.csv` data set provides more data about the resources required for each project. Each line in this file represents a resource required by a project:

Feature	Description
id	A <code>project_id</code> value from the <code>train.csv</code> file. Example: p036502
description	Description of the resource. Example: Tenor Saxophone Reeds, Box of 25
quantity	Quantity of the resource required. Example: 3
price	Price of the resource required. Example: 9.95

Note: Many projects require multiple resources. The `id` value corresponds to a `project_id` in `train.csv`, so you use it as a key to retrieve all resources needed for a project:

The data set contains the following label (the value you will attempt to predict):

Label	Description
<code>project_is_approved</code>	A binary flag indicating whether DonorsChoose approved the project. A value of <code>0</code> indicates the project was not approved, and a value of <code>1</code> indicates the project was approved.



Notes on the Essay Data

Prior to May 17, 2016, the prompts for the essays were as follows:

- **project_essay_1:** "Introduce us to your classroom"
- **project_essay_2:** "Tell us more about your students"
- **project_essay_3:** "Describe how your students will use the materials you're requesting"
- **project_essay_3:** "Close by sharing why your project will make a difference"

Starting on May 17, 2016, the number of essays was reduced from 4 to 2, and the prompts for the first 2 essays were changed to the following:

- **project_essay_1:** "Describe your students: What makes your students special? Specific details about their background, your neighborhood, and your school are all helpful."
- **project_essay_2:** "About your project: How will these materials make a difference in your students' learning and improve their school lives?"

For all projects with `project_submitted_datetime` of 2016-05-17 and later, the values of `project_essay_3` and `project_essay_4` will be NaN.

```
In [196]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

from chart_studio import plotly
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
from collections import Counter
```

1.1 Reading Data

```
In [197]: project_data = pd.read_csv('train_data.csv')
resource_data = pd.read_csv('resources.csv')
```

```
In [198]: project_data.shape
```

```
Out[198]: (109248, 17)
```

```
In [199]: print("Number of data points in train data", project_data.shape)
print('-'*50)
print("The attributes of data :", project_data.columns.values)
project_data.project_is_approved.value_counts()
```

Number of data points in train data (109248, 17)

```
-----
The attributes of data : ['Unnamed: 0' 'id' 'teacher_id' 'teacher_prefix' 'school_state'
'project_submitted_datetime' 'project_grade_category'
'project_subject_categories' 'project_subject_subcategories'
'project_title' 'project_essay_1' 'project_essay_2' 'project_essay_3'
'project_essay_4' 'project_resource_summary'
'teacher_number_of_previously_posted_projects' 'project_is_approved']
```

```
Out[199]: 1    92706
0    16542
Name: project_is_approved, dtype: int64
```

```
In [200]: print("Number of data points in train data", resource_data.shape)
print(resource_data.columns.values)
resource_data.head(2)
```

Number of data points in train data (1541272, 4)
['id' 'description' 'quantity' 'price']

```
Out[200]:
```

	id	description	quantity	price
0	p233245	LC652 - Lakeshore Double-Space Mobile Drying Rack	1	149.00
1	p069063	Bouncy Bands for Desks (Blue support pipes)	3	14.95

1.2 preprocessing of project_subject_categories

```

In [201]: categories = list(project_data['project_subject_categories'].values)
# remove special characters from list of strings python: https://stackoverflow.com/

# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python

cat_list = []
for i in categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split(): # this will split each of the category based on space
            j=j.replace('The', '') # if we have the words "The" we are going to remove them
        j = j.replace(' ', '') # we are replacing all the ' ' (space) with '' (empty)
        temp+=j.strip()+" " # " abc ".strip() will return "abc", remove the trailing spaces
    temp = temp.replace('&', '_') # we are replacing the & value into _
    cat_list.append(temp.strip())

project_data['clean_categories'] = cat_list
project_data.drop(['project_subject_categories'], axis=1, inplace=True)

from collections import Counter
my_counter = Counter()
for word in project_data['clean_categories'].values:
    my_counter.update(word.split())

cat_dict = dict(my_counter)
sorted_cat_dict = dict(sorted(cat_dict.items(), key=lambda kv: kv[1]))

```

1.3 preprocessing of project_subject_subcategories

```

In [202]: sub_categories = list(project_data['project_subject_subcategories'].values)
# remove special characters from list of strings python: https://stackoverflow.com/

# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python

sub_cat_list = []
for i in sub_categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split(): # this will split each of the category based on space
            j=j.replace('The', '') # if we have the words "The" we are going to remove them
        j = j.replace(' ', '') # we are replacing all the ' ' (space) with '' (empty)
        temp +=j.strip()+" #" "abc ".strip() will return "abc", remove the trailing spaces
        temp = temp.replace('&','_')
    sub_cat_list.append(temp.strip())

project_data['clean_subcategories'] = sub_cat_list
project_data.drop(['project_subject_subcategories'], axis=1, inplace=True)

# count of all the words in corpus python: https://stackoverflow.com/a/22898595/4090899
my_counter = Counter()
for word in project_data['clean_subcategories'].values:
    my_counter.update(word.split())

sub_cat_dict = dict(my_counter)
sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))

```



```
In [206]: print(preprocessed_project_grade_categories[12])
print("="*50)
print(preprocessed_project_grade_categories[502])
print("="*50)
print(preprocessed_project_grade_categories[5002])
print("="*50)
print(preprocessed_project_grade_categories[50002])
print("="*50)
print(preprocessed_project_grade_categories[100013])
print("="*50)
```

```
6_8
=====
6_8
=====
6_8
=====
3_5
=====
PreK_2
=====
```

1.5 preprocessing of teacher prefix

```
In [207]: project_data['teacher_prefix'].value_counts()
```

```
Out[207]: Mrs.      57269
Ms.       38955
Mr.       10648
Teacher   2360
Dr.        13
Name: teacher_prefix, dtype: int64
```

Replacing Nan Values with maximum frequencies values i.e Mrs.

```
In [208]: project_data['teacher_prefix'] = project_data['teacher_prefix'].fillna('Mrs.')
```

```
In [209]: def replace_cate(lst):          # Removing (.) in Mrs.
            return lst.replace('.', '')
project_data['teacher_prefix'] = project_data['teacher_prefix'].astype(str).apply(
```

```
In [210]: preprocessed_teacher_prefix = []
for teach_prefix in tqdm(project_data["teacher_prefix"]):
    preprocessed_teacher_prefix.append(teach_prefix.strip())
```

```
100%|████████████████████████████████████████████████████████████████████████████████| 109
248/109248 [00:00<00:00, 2235597.90it/s]
```

```
In [211]: print(preprocessed_teacher_prefix[1])
print("="*50)
print(preprocessed_teacher_prefix[50])
print("="*50)
project_data.teacher_prefix.value_counts()
```

```
Mr
=====
Mrs
=====
```

```
Out[211]: Mrs      57272
Ms       38955
Mr       10648
Teacher   2360
Dr         13
Name: teacher_prefix, dtype: int64
```

1.6 Adding a new feature Number of words in title(optional)

```
In [212]: title_word_count = []
```


```
In [213]: for a in project_data["project_title"] :
b = len(a.split())
title_word_count.append(b)
```

```
In [214]: project_data["title_word_count"] = title_word_count
```

In [215]: `project_data.head(5)`

Out[215]:

	Unnamed: 0	id	teacher_id	teacher_prefix	school_state	project_sul
0	160221	p253737	c90749f5d961ff158d4b4d1e7dc665fc	Mrs	IN	20
1	140945	p258326	897464ce9ddc600bced1151f324dd63a	Mr	FL	20
2	21895	p182444	3465aaf82da834c0582ebd0ef8040ca0	Ms	AZ	20
3	45	p246581	f3cb9bffbba169bef1a77b243e620b60	Mrs	KY	20
4	172407	p104768	be1f7507a41f8479dc06f047086a39ec	Mrs	TX	20



combining 4 essays into 1 essay

```
In [216]: # merge two column text dataframe:
project_data["essay"] = project_data["project_essay_1"].map(str) + \
    project_data["project_essay_2"].map(str) + \
    project_data["project_essay_3"].map(str) + \
    project_data["project_essay_4"].map(str)
```

Adding a new feature Number of words in essay

```
In [217]: essay_word_count=[]
```

```
In [218]: for ess in project_data["essay"] :  
          c = len(ess.split())  
          essay_word_count.append(c)
```

```
In [219]: project_data["essay_word_count"] = essay_word_count
```

```
In [220]: project_data.head(2)
```

Out[220]:

Unnamed: 0	id	teacher_id	teacher_prefix	school_state	project_sul
0	160221 p253737	c90749f5d961ff158d4b4d1e7dc665fc	Mrs	IN	20
1	140945 p258326	897464ce9ddc600bced1151f324dd63a	Mr	FL	20

```
In [221]: project_data.head(2)
```

Out[221]:

Unnamed: 0	id	teacher_id	teacher_prefix	school_state	project_sul
0	160221 p253737	c90749f5d961ff158d4b4d1e7dc665fc	Mrs	IN	20
1	140945 p258326	897464ce9ddc600bced1151f324dd63a	Mr	FL	20

Train Test split

```
In [222]: # train test split using sklearn.model selection
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(project_data, project_data['p
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.33,
```

```
In [223]: project_data.head()
```

Out[223]:


	Unnamed: 0	id	teacher_id	teacher_prefix	school_state	project_sul
0	160221	p253737	c90749f5d961ff158d4b4d1e7dc665fc	Mrs	IN	20
1	140945	p258326	897464ce9ddc600bced1151f324dd63a	Mr	FL	20
2	21895	p182444	3465aaf82da834c0582ebd0ef8040ca0	Ms	AZ	20
3	45	p246581	f3cb9bffbba169bef1a77b243e620b60	Mrs	KY	20
4	172407	p104768	be1f7507a41f8479dc06f047086a39ec	Mrs	TX	20

```
In [224]: X_train.drop(['project_is_approved'], axis=1, inplace=True)
X_test.drop(['project_is_approved'], axis=1, inplace=True)
X_cv.drop(['project_is_approved'], axis=1, inplace=True)
```

```
In [225]: X_train.head(2)
```

Out[225]:

	Unnamed: 0	id	teacher_id	teacher_prefix	school_state	projec
21389	75072	p146651	c50bfe62d50fdb9dc30d5bfa107e4c5a	Ms	PA	
34546	74550	p090750	ae699781d1e0eb2fd89d02008ec23b05	Mrs	IA	



1.8 Text preprocessing

```
In [226]: # printing some random reviews

print(X_train['essay'].values[0])
print("="*50)
print(X_train['essay'].values[500])
print("="*50)
```

As a teacher in a low-income/high poverty school, my students are faced with several challenges both inside and outside the classroom. Despite the many challenges they face, I want to provide an environment that is safe, fun and engaging so they can get the most out of their day in my room.\r\n\r\nMy students are creative, clever, and very unique. They like to move, to help and love lots of positive attention. Many of them are being raised in single-parent households and receive a free lunch based on their socioeconomic status. These things may prevent them from getting ahead early in life and may not provide them with the life experiences many of us see as \"typical\". From the minute they walk in the door of my classroom, I focus on their potential and growth while they are with me. I may not be able to control their home lives but I can certainly control their experience during the school day. By doing this in a creative and positive way, I am hopeful to inspire my students to continue on a path of academic excellence.\r\n\r\nBeing able to display their hard work gives my students a sense of pride and accomplishment. Their face shows it when their hard work is hung in the hallway and in the classroom. They have satisfaction when they can bring home their work to their families.\r\n\r\n\r\nMy students work hard and love when they are able to share their work with others. We are requesting some copy paper and toner for our printer so that we can continue to share our work with others. These much needed supplies will allow our creativity and passion to continue throughout the school year.nannan

=====
As a teacher in a low-income/high poverty school district, my students are faced with several challenges both in and out of the classroom. When the students enter the class each morning they know they are safe. Our classroom is home and we are family. The students know this is a place they can let their guards down and be their unique selves. They know they are safe to take risks and engage in challenges to grow in academics and social and emotional skills. My babies know I am their \"Mama Bear\" for five days a week.This donation will be used daily in small group math interactive activities. This donation will also be used for students to read comfortably books of their choice. It will also create an area where the students will feel comfortable to take risks and increase their daily education. This will improve their lives because they will increase education through peer discussions and interactions more frequently. \r\n\r\n\r\nA new rug helps to create a classroom where students feel free to learn in a setting of their choice. Students can move away from the standard desks and into learning centers where they have more opportunities to have academic discussions.nannan

=====

In [227]: [# https://stackoverflow.com/a/47091490/4084039](https://stackoverflow.com/a/47091490/4084039)

```
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase
```

```
In [228]: sent = decontracted(X_train['essay'].values[20000])
print(sent)
print("="*50)
```

My students are fun, hard-working and excited about learning. It is not fun every minute of every day but we do a pretty good job of making learning fun. They enter the room with the confidence and enthusiasm for learning. They travel the world through books and make real world connections through talk and interactions. We always talk about the importance of our education so they always strive to do their very best. Their bright faces are my motivation to work harder. Before the historic flooding in the city of Baton Rouge, I welcomed my students back to school with a smiling face and warm hugs. Many came to school excited but with minimal supplies. We were in school two days before the flood. Our school is neighborhood, which many of our students live, took in water. A lot of students lost everything including school supplies they had not brought to school. The project will furnish my students with the supplies they need and take the added stress off of the parents who are already dealing with a lot of devastation. Thanks a bunch!nannan

=====


```
In [229]: # \r \n \t remove from string python: http://texthandler.com/info/remove-line-breaks
sent = sent.replace('\r', ' ')
sent = sent.replace('\n', ' ')
sent = sent.replace('\t', ' ')
print(sent)
```

My students are fun, hard-working and excited about learning. It is not fun every minute of every day but we do a pretty good job of making learning fun. They enter the room with the confidence and enthusiasm for learning. They travel the world through books and make real world connections through talk and interactions. We always talk about the importance of our education so they always strive to do their very best. Their bright faces are my motivation to work harder. Before the historic flooding in the city of Baton Rouge, I welcomed my students back to school with a smiling face and warm hugs. Many came to school excited but with minimal supplies. We were in school two days before the flood. Our school is neighborhood, which many of our students live, took in water. A lot of students lost everything including school supplies they had not brought to school. The project will furnish my students with the supplies they need and take the added stress off of the parents who are already dealing with a lot of devastation. Thanks a bunch!nannan

```
In [230]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
print(sent)
```

My students are fun hard working and excited about learning It is not fun every minute of every day but we do a pretty good job of making learning fun They enter the room with the confidence and enthusiasm for learning They travel the world through books and make real world connections through talk and interactions We always talk about the importance of our education so they always strive to do their very best Their bright faces are my motivation to work harder Before the historic flooding in the city of Baton Rouge I welcomed my students back to school with a smiling face and warm hugs Many came to school excited but with minimal supplies We were in school two days before the flood Our school is neighborhood which many of our students live took in water A lot of students lost everything including school supplies they had not brought to school The project will furnish my students with the supplies they need and take the added stress off of the parents who are already dealing with a lot of devastation Thanks a bunch nannan


```
In [235]: preprocessed_essays_test
```

oked would not know pride students show take immaculate care instruments equipment represent school pride everything funded materials requested placed band hall new school year main component binders page protectors binder home base beginner band help teach organizational skills not help band program studies show skills transfer areas education lives binder store notes practice logs handouts importantly music dividers allow students insert music make marks throughout removed no longer need dividers assist moving quickly throughout binders rehearsal starting beginning band sort organization recipe success pencils markers erasers paper help keep supplies students hands student purchase many supplies classrooms difficult every student bring supplies music classroom helpful materials specifically music create learn ease nannan',

'title one school 70 students free reduced priced lunch number may people think hear school not hear students working together everyday solve problems within classroom despite students adversities eager come school learn day students hardest working students able work happiest students despite everything go daily basis even make school students coming traumatic backgrounds school safe place students come educators students families want able give best materials help foster learning goal give students materials need successful students love active learn always wanting move lessons want students understand important fit healthy attentive learners asked students needed classroom next year

1.8.3 Preprocessed cross validation data

```
In [236]: preprocessed_essays_cv=preprocess_textual(X_cv['essay'].values)
```

```
100%|██████████████████████████████████████████|  
24155/24155 [00:14<00:00, 1699.17it/s]
```

1.9 preprocessing of project title

```
In [237]: # printing some random project titles.
print(project_data['project_title'].values[1])
print("="*50)
print(project_data['project_title'].values[1501])
print("="*50)
print(project_data['project_title'].values[10001])
print("="*50)
print(project_data['project_title'].values[20001])
print("="*50)
```

Wanted: Projector for Hungry Learners
=====

Making Every Day at School Count
=====

Becoming 'Readerly' Readers
=====

The Beautiful Life of a Butterfly
=====

```
In [238]: title = decontracted(X_train['project_title'].values[2000])
```

1.9.1 Preprocessing of Project Title(Train)

```
In [239]: preprocessed_titles_train=preprocess_textual(X_train["project_title"].values)
```

```
100%|████████████████████████████████████████████████████████████████████████████████|
49041/49041 [00:01<00:00, 44925.38it/s]
```

1.9.2 Preprocessing of Project Title(Test)

```
In [240]: preprocessed_titles_test=preprocess_textual(X_test["project_title"])
```

```
100%|████████████████████████████████████████████████████████████████████████████████|
36052/36052 [00:00<00:00, 43200.58it/s]
```

```
In [241]: preprocessed_titles_test[10]
```

```
Out[241]: 'chromebooks empower students 21st century learning'
```

1.9.2 Preprocessing of Project Title(CV)

```
In [242]: preprocessed_titles_cv=preprocess_textual(X_cv["project_title"])
```

```
100%|████████████████████████████████████████████████████████████████████████████████|
24155/24155 [00:00<00:00, 46958.40it/s]
```

```
In [243]: preprocessed_titles_cv[10]
```

```
Out[243]: 'newbery books 5th grade'
```

1.5 Preparing data for models

```
In [244]: project_data.columns
```

```
Out[244]: Index(['Unnamed: 0', 'id', 'teacher_id', 'teacher_prefix', 'school_state',
               'project_submitted_datetime', 'project_grade_category', 'project_title',
               'project_essay_1', 'project_essay_2', 'project_essay_3',
               'project_essay_4', 'project_resource_summary',
               'teacher_number_of_previously_posted_projects', 'project_is_approved',
               'clean_categories', 'clean_subcategories', 'title_word_count', 'essay',
               'essay_word_count'],
              dtype='object')
```

we are going to consider

- school_state : categorical data
- clean_categories : categorical data
- clean_subcategories : categorical data
- project_grade_category : categorical data
- teacher_prefix : categorical data

- project_title : text data
- text : text data
- project_resource_summary: text data (optinal)

- quantity : numerical (optinal)
- teacher_number_of_previously_posted_projects : numerical
- price : numerical

```
In [245]: project_data.head(10)
```

Out[245]:

	Unnamed: 0	id	teacher_id	teacher_prefix	school_state	project_su
0	160221	p253737	c90749f5d961ff158d4b4d1e7dc665fc	Mrs	IN	20
1	140945	p258326	897464ce9ddc600bcd1151f324dd63a	Mr	FL	20
2	21895	p182444	3465aaf82da834c0582ebd0ef8040ca0	Ms	AZ	20
3	45	p246581	f3cb9bffbb169bef1a77b243e620b60	Mrs	KY	20
4	172407	p104768	be1f7507a41f8479dc06f047086a39ec	Mrs	TX	20
5	141660	p154343	a50a390e8327a95b77b9e495b58b9a6e	Mrs	FL	20
6	21147	p099819	9b40170bfa65e399981717ee8731efc3	Mrs	CT	20
7	94142	p092424	5bfd3d12fae3d2fe88684bbac570c9d2	Ms	GA	20
8	112489	p045029	487448f5226005d08d36bdd75f095b31	Mrs	SC	20

Unnamed: 0		id	teacher_id	teacher_prefix	school_state	project_su
9	158561	p001713	140eeac1885c820ad5592a409a3a8994	Ms	NC	21

1.5.1 Vectorizing Categorical data

In [246]: `print(project_data.shape)`

(109248, 20)

one hot vector for clean categories of Projects (train,test,cv)

```
In [247]: # we use count vectorizer to convert the values into one hot vectors

from sklearn.feature_extraction.text import CountVectorizer

vectorizer_proj = CountVectorizer(min_df=10,ngram_range=(1,1), max_features=5000)
vectorizer_proj.fit(X_train['clean_categories'].values)

categories_one_hot_train = vectorizer_proj.transform(X_train['clean_categories'].values)
categories_one_hot_test = vectorizer_proj.transform(X_test['clean_categories'].values)
categories_one_hot_cv = vectorizer_proj.transform(X_cv['clean_categories'].values)

print(vectorizer_proj.get_feature_names())

print("Shape of matrix of Train data after one hot encoding ",categories_one_hot_train.shape)
print("Shape of matrix of Test data after one hot encoding ",categories_one_hot_test.shape)
print("Shape of matrix of CV data after one hot encoding ",categories_one_hot_cv.shape)

['appliedlearning', 'care_hunger', 'health_sports', 'history_civics', 'literacy_language', 'math_science', 'music_arts', 'specialneeds', 'warmth']
Shape of matrix of Train data after one hot encoding (49041, 9)
Shape of matrix of Test data after one hot encoding (36052, 9)
Shape of matrix of CV data after one hot encoding (24155, 9)
```

one hot vector for clean subcategories (train,test,cv)

```
In [248]: # we use count vectorizer to convert the values into one

vectorizer_sub_proj = CountVectorizer(min_df=10,ngram_range=(1,1), max_features=50)
vectorizer_sub_proj.fit(X_train['clean_subcategories'].values)

sub_categories_one_hot_train = vectorizer_sub_proj.transform(X_train['clean_subcategories'])
sub_categories_one_hot_test = vectorizer_sub_proj.transform(X_test['clean_subcategories'])
sub_categories_one_hot_cv = vectorizer_sub_proj.transform(X_cv['clean_subcategories'])

print(vectorizer_sub_proj.get_feature_names())

print("Shape of matrix of Train data after one hot encoding ",sub_categories_one_hot_train.shape)
print("Shape of matrix of Test data after one hot encoding ",sub_categories_one_hot_test.shape)
print("Shape of matrix of Cross Validation data after one hot encoding ",sub_categories_one_hot_cv.shape)
```

```
['appliedsciences', 'care_hunger', 'charactereducation', 'civics_government',
 'college_careerprep', 'communityservice', 'earlydevelopment', 'economics', 'environmentalscience', 'esl', 'extracurricular', 'financialliteracy', 'foreignlanguages', 'gym_fitness', 'health_lifescience', 'health_wellness', 'history_geography', 'literacy', 'literature_writing', 'mathematics', 'music', 'nutritioneducation', 'other', 'parentinvolvement', 'performingarts', 'socialsciences', 'specialneeds', 'teamsports', 'visualarts', 'warmth']
Shape of matrix of Train data after one hot encoding (49041, 30)
Shape of matrix of Test data after one hot encoding (36052, 30)
Shape of matrix of Cross Validation data after one hot encoding (24155, 30)
```

```
In [249]: sub_categories_one_hot_cv[0].toarray()
```

```
Out[249]: array([[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0]], dtype=int64)
```

One hot vector for school states(train,test,cv)

```
In [250]: my_counter = Counter()
for state in project_data['school_state'].values:
    my_counter.update(state.split())
```

```
In [251]: school_state_cat_dict = dict(my_counter)
sorted_school_state_cat_dict = dict(sorted(school_state_cat_dict.items(), key=lambda item: item[0]))
```



```
In [252]: ## Using count vectorizer to convert the values into one hot encoded features

vectorizer_states = CountVectorizer(min_df=10,ngram_range=(1,1), max_features=5000)
vectorizer_states.fit(X_train['school_state'].values)

school_state_categories_one_hot_train = vectorizer_states.transform(X_train['school_state'])
school_state_categories_one_hot_test = vectorizer_states.transform(X_test['school_state'])
school_state_categories_one_hot_cv = vectorizer_states.transform(X_cv['school_state'])

print(vectorizer_states.get_feature_names())

print("Shape of matrix of Train data after one hot encoding ",school_state_categories_one_hot_train.shape)
print("Shape of matrix of Test data after one hot encoding ",school_state_categories_one_hot_test.shape)
print("Shape of matrix of Cross Validation data after one hot encoding ",school_state_categories_one_hot_cv.shape)

['ak', 'al', 'ar', 'az', 'ca', 'co', 'ct', 'dc', 'de', 'fl', 'ga', 'hi', 'ia', 'id', 'il', 'in', 'ks', 'ky', 'la', 'ma', 'md', 'me', 'mi', 'mn', 'mo', 'ms', 'mt', 'nc', 'nd', 'ne', 'nh', 'nj', 'nm', 'nv', 'ny', 'oh', 'ok', 'or', 'pa', 'ri', 'sc', 'sd', 'tn', 'tx', 'ut', 'va', 'vt', 'wa', 'wi', 'wv', 'wy']
Shape of matrix of Train data after one hot encoding (49041, 51)
Shape of matrix of Test data after one hot encoding (36052, 51)
Shape of matrix of Cross Validation data after one hot encoding (24155, 51)
```

one hot vector for Project grade category (train,test,cv)

```
In [253]: my_counter = Counter()
for project_grade in preprocessed_project_grade_categories:
    my_counter.update(project_grade.split())
```

```
In [254]: project_grade_cat_dict = dict(my_counter)
sorted_project_grade_cat_dict = dict(sorted(project_grade_cat_dict.items(), key=lambda item: item[1], reverse=True))
```

```
In [255]: ## we use count vectorizer to convert the values into one hot encoded features
vectorizer_grade = CountVectorizer(min_df=10,ngram_range=(1,1), max_features=5000)
vectorizer_grade.fit(X_train['project_grade_category'].values)

project_grade_categories_one_hot_train = vectorizer_grade.transform(X_train['project_grade_category'])
project_grade_categories_one_hot_test = vectorizer_grade.transform(X_test['project_grade_category'])
project_grade_categories_one_hot_cv = vectorizer_grade.transform(X_cv['project_grade_category'])

print(vectorizer_grade.get_feature_names())

print("Shape of matrix of Train data after one hot encoding ",project_grade_categories_one_hot_train.shape)
print("Shape of matrix of Test data after one hot encoding ",project_grade_categories_one_hot_test.shape)
print("Shape of matrix of Cross Validation data after one hot encoding ",project_grade_categories_one_hot_cv.shape)

['12', 'grades', 'prek']
Shape of matrix of Train data after one hot encoding (49041, 3)
Shape of matrix of Test data after one hot encoding (36052, 3)
Shape of matrix of Cross Validation data after one hot encoding (24155, 3)
```

One hot vector for teacher prefix(train,test,cv)

```
In [256]: vectorizer_teacher = CountVectorizer()
vectorizer_teacher.fit(X_train['teacher_prefix'].values) # fit has to happen only

# we use the fitted CountVectorizer to convert the text to vector
teacher_prefix_categories_one_hot_train = vectorizer_teacher.transform(X_train['teacher_prefix'])
teacher_prefix_categories_one_hot_cv = vectorizer_teacher.transform(X_cv['teacher_prefix'])
teacher_prefix_categories_one_hot_test = vectorizer_teacher.transform(X_test['teacher_prefix'])

print("After vectorizations")
print("Shape of matrix of Train data after one hot encoding",teacher_prefix_categories_one_hot_train.shape)
print("Shape of matrix of cv data after one hot encoding",teacher_prefix_categories_one_hot_cv.shape)
print("Shape of matrix of Test data after one hot encoding",teacher_prefix_categories_one_hot_test.shape)
print(vectorizer_teacher.get_feature_names())
print("=="*100)
```

After vectorizations

Shape of matrix of Train data after one hot encoding (49041, 5) (49041,)

Shape of matrix of cv data after one hot encoding (24155, 5) (24155,)

Shape of matrix of Test data after one hot encoding (36052, 5) (36052,)

['dr', 'mr', 'mrs', 'ms', 'teacher']

=====

1.11 Vectorizing text data

A) Bag of words

BOW train data essays

```
In [257]: # We are considering only the words which appeared in at least 10 documents(rows)
vectorizer_bow_essay = CountVectorizer(min_df=10,max_features=5000) #selecting top 5000 words
vectorizer_bow_essay.fit(preprocessed_essays_train)

text_bow_train = vectorizer_bow_essay.transform(preprocessed_essays_train)

print("Shape of matrix after one hot encoding ",text_bow_train.shape)
```

Shape of matrix after one hot encoding (49041, 5000)

BOW (test essays)

```
In [258]: text_bow_test = vectorizer_bow_essay.transform(preprocessed_essays_test)
print("Shape of matrix after one hot encoding ",text_bow_test.shape)
```

Shape of matrix after one hot encoding (36052, 5000)

Bow (cv essays)

```
In [259]: text_bow_cv = vectorizer_bow_essay.transform(preprocessed_essays_cv)
print("Shape of matrix after one hot encoding ",text_bow_cv.shape)
```

Shape of matrix after one hot encoding (24155, 5000)

Bow(train titles)

```
In [260]: vectorizer_bow_title = CountVectorizer(min_df=10)
vectorizer_bow_title.fit(preprocessed_titles_train)
title_bow_train = vectorizer_bow_title.transform(preprocessed_titles_train)
print("Shape of matrix after one hot encoding ",title_bow_train.shape)
```

Shape of matrix after one hot encoding (49041, 1990)

Bow(test titles)

```
In [261]: title_bow_test = vectorizer_bow_title.transform(preprocessed_titles_test)
print("Shape of matrix after one hot encoding ",title_bow_test.shape)
```

Shape of matrix after one hot encoding (36052, 1990)

Bow(cv titles)

```
In [262]: title_bow_cv = vectorizer_bow_title.transform(preprocessed_titles_cv)
print("Shape of matrix after one hot encoding ",title_bow_cv.shape)
```

Shape of matrix after one hot encoding (24155, 1990)

B) Tfidf

tfidf(train essays)

```
In [263]: from sklearn.feature_extraction.text import TfidfVectorizer

vectorizer_tfidf_essay = TfidfVectorizer(min_df=10,max_features=5000) #Considering
vectorizer_tfidf_essay.fit(preprocessed_essays_train)

text_tfidf_train = vectorizer_tfidf_essay.transform(preprocessed_essays_train)
print("Shape of matrix after one hot encoding ",text_tfidf_train.shape)
```

Shape of matrix after one hot encoding (49041, 5000)

tfidf(test essays)

```
In [264]: text_tfidf_test = vectorizer_tfidf_essay.transform(preprocessed_essays_test)
print("Shape of matrix after one hot encoding ",text_tfidf_test.shape)
```

Shape of matrix after one hot encoding (36052, 5000)

tfidf(cv essays)

```
In [265]: text_tfidf_cv = vectorizer_tfidf_essay.transform(preprocessed_essays_cv)
print("Shape of matrix after one hot encoding ",text_tfidf_cv.shape)
```

Shape of matrix after one hot encoding (24155, 5000)

tfidf(train titles)

```
In [266]: vectorizer_tfidf_titles = TfidfVectorizer(min_df=10)

vectorizer_tfidf_titles.fit(preprocessed_titles_train)
title_tfidf_train = vectorizer_tfidf_titles.transform(preprocessed_titles_train)
print("Shape of matrix after one hot encoding ",title_tfidf_train.shape)
```

Shape of matrix after one hot encoding (49041, 1990)

tfidf(test titles)

```
In [267]: title_tfidf_test = vectorizer_tfidf_titles.transform(preprocessed_titles_test)
print("Shape of matrix after one hot encoding ",title_tfidf_test.shape)
```

Shape of matrix after one hot encoding (36052, 1990)

tfidf(cv titles)

```
In [268]: title_tfidf_cv = vectorizer_tfidf_titles.transform(preprocessed_titles_cv)
print("Shape of matrix after one hot encoding ",title_tfidf_cv.shape)
```

Shape of matrix after one hot encoding (24155, 1990)

1.5.2.3 Using Pretrained Models: Avg W2V

```
In [269]: # stronging variables into pickle files python: http://www.jessicayung.com/how-to
# make sure you have the glove_vectors file
with open('glove_vectors', 'rb') as f:
    model = pickle.load(f)
    glove_words = set(model.keys())
```

train essays

```
In [270]: def avg_word_vec(preprocessed_data):
    avg_w2v_vectors = [];
    for sentence in tqdm(preprocessed_data): # for each review/sentence
        vector = np.zeros(300) # as word vectors are of zero length
        cnt_words = 0; # num of words with a valid vector in the sentence/review
        for word in sentence.split(): # for each word in a review/sentence
            if word in glove_words:
                vector += model[word]
                cnt_words += 1
        if cnt_words != 0:
            vector /= cnt_words
        avg_w2v_vectors.append(vector)
    return avg_w2v_vectors
```

```
In [271]: avg_w2v_vectors_train = avg_word_vec(preprocessed_essays_train)
print(len(avg_w2v_vectors_train))
print(len(avg_w2v_vectors_train[0]))
```

```
100%|████████████████████████████████████████████████████████████████████████████████|
| 49041/49041 [00:11<00:00, 4251.97it/s]
```

```
49041
300
```

test essays

```
In [272]: avg_w2v_vectors_test = avg_word_vec(preprocessed_essays_test)
print(len(avg_w2v_vectors_test))
print(len(avg_w2v_vectors_test[0]))
```

```
100%|████████████████████████████████████████████████████████████████████████████████|
| 36052/36052 [00:08<00:00, 4220.36it/s]
```

```
36052
300
```

cv essays

```
In [273]: avg_w2v_vectors_cv = avg_word_vec(preprocessed_essays_cv)
print(len(avg_w2v_vectors_cv))
print(len(avg_w2v_vectors_cv[0]))
```

```
100%|████████████████████████████████████████████████████████████████████████████████|
| 24155/24155 [00:06<00:00, 3758.37it/s]
```

```
24155
300
```

train titles

```
In [274]: avg_w2v_vectors_titles_train = avg_word_vec(preprocessed_titles_train)
print(len(avg_w2v_vectors_titles_train))
print(len(avg_w2v_vectors_titles_train[0]))
```

```
100%|████████████████████████████████████████████████████████████████████████████████|
49041/49041 [00:00<00:00, 79370.90it/s]
```

```
49041
300
```

test titles

```
In [275]: avg_w2v_vectors_titles_test = avg_word_vec(preprocessed_titles_test)
print(len(avg_w2v_vectors_titles_test))
print(len(avg_w2v_vectors_titles_test[0]))
```

```
100%|████████████████████████████████████████████████████████████████████████████████|
36052/36052 [00:00<00:00, 46763.63it/s]
```

```
36052
300
```

Cv titles

```
In [276]: avg_w2v_vectors_titles_cv = avg_word_vec(preprocessed_titles_cv)
print(len(avg_w2v_vectors_titles_cv))
print(len(avg_w2v_vectors_titles_cv[0]))
```

```
100%|████████████████████████████████████████████████████████████████████████████████|
24155/24155 [00:00<00:00, 98446.65it/s]
```

```
24155
300
```

using pretrained models : Tfidf weighted W2V

train essays

```
In [277]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
tfidf_model = TfidfVectorizer()
tfidf_model.fit(preprocessed_essays_train)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(tfidf_model.get_feature_names(), list(tfidf_model.idf_)))
tfidf_words = set(tfidf_model.get_feature_names())
```

```
In [278]: def tfidfWV(preprocessed_data):
# average Word2Vec
# compute average word2vec for each review.
tfidf_w2v_vectors = []; # the avg-w2v for each sentence/review is stored in t
for sentence in tqdm(preprocessed_data): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split()))
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    tfidf_w2v_vectors.append(vector)
print("The length of TFIDF word to vec is ", len(tfidf_w2v_vectors))
print("The length of TFIDF word to vec of index 0 is ", len(tfidf_w2v_vectors[0]))
return tfidf_w2v_vectors
```

```
In [279]: tfidf_w2v_vectors_train=tfidfWV(preprocessed_essays_train)
```

```
100%|████████████████████████████████████████████████████████████████████████████████|
49041/49041 [01:24<00:00, 577.39it/s]
```

The length of TFIDF word to vec is 49041

The length of TFIDF word to vec of index 0 is 300

Test essays

```
In [280]: tfidf_w2v_vectors_test=tfidfWV(preprocessed_essays_test)
```

```
100%|████████████████████████████████████████████████████████████████████████████████|
36052/36052 [01:00<00:00, 599.19it/s]
```

The length of TFIDF word to vec is 36052

The length of TFIDF word to vec of index 0 is 300

cv essays

```
In [281]: tfidf_w2v_vectors_cv = tfidfWV(preprocessed_essays_cv)
```

```
100%|████████████████████████████████████████████████████████████████████████████████|
24155/24155 [00:40<00:00, 600.07it/s]
```

The length of TFIDF word to vec is 24155

The length of TFIDF word to vec of index 0 is 300

train titles

- 4.Title word Count (introduced by us)
- 5.Essay word Count (introduced by us)

1) Price

In [286]:

```
# https://stackoverflow.com/questions/22407798/how-to-reset-a-dataframes-indexes-j
price_data = resource_data.groupby('id').agg({'price':'sum', 'quantity':'sum'}).reset_index()
price_data.head(4)
```

Out[286]:

	id	price	quantity
0	p000001	459.56	7
1	p000002	515.89	21
2	p000003	298.97	4
3	p000004	1113.69	98

In [287]:

```
# join two dataframes in python:
X_train = pd.merge(X_train, price_data, on='id', how='left')
X_test = pd.merge(X_test, price_data, on='id', how='left')
X_cv = pd.merge(X_cv, price_data, on='id', how='left')
```

In [288]:

X_train

49037	103804	p217443	a4adec7a1d197fcb38cd8d04c55736a9	Mrs	CA
49038	83859	p233438	cb0ff2b4f2830f05248d423aad50ffe3	Ms	ID
49039	8111	p202273	7ac85d702a9582190b2c068c33f5a65f	Ms	MI
49040	32480	p258261	95892c1b13f2aafce2643caf0d1f5f77	Ms	GA

49041 rows × 21 columns

```
In [289]: from sklearn.preprocessing import Normalizer

normalizer = Normalizer()

# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.

normalizer.fit(X_train['price'].values.reshape(1, -1))

price_train = normalizer.transform(X_train['price'].values.reshape(1, -1))
price_cv = normalizer.transform(X_cv['price'].values.reshape(1, -1))
price_test = normalizer.transform(X_test['price'].values.reshape(1, -1))

print("After vectorizations")
print(price_train.shape, y_train.shape)
print(price_cv.shape, y_cv.shape)
print(price_test.shape, y_test.shape)
print("=="*100)
```

After vectorizations

(1, 49041) (49041,)

(1, 24155) (24155,)

(1, 36052) (36052,)

=====
=====

2) Quantity

```
In [290]: normalizer = Normalizer()

# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.

normalizer.fit(X_train['quantity'].values.reshape(1, -1))

quantity_train = normalizer.transform(X_train['quantity'].values.reshape(1, -1))
quantity_cv = normalizer.transform(X_cv['quantity'].values.reshape(1, -1))
quantity_test = normalizer.transform(X_test['quantity'].values.reshape(1, -1))

print("After vectorizations")
print(quantity_train.shape, y_train.shape)
print(quantity_cv.shape, y_cv.shape)
print(quantity_test.shape, y_test.shape)
print("=="*100)
```

After vectorizations

(1, 49041) (49041,)

(1, 24155) (24155,)

(1, 36052) (36052,)

=====
=====

3) Number of Projects previously proposed by Teacher

```
In [291]: normalizer = Normalizer()

# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.

normalizer.fit(X_train['teacher_number_of_previously_posted_projects'].values.reshape(1, -1))

prev_projects_train = normalizer.transform(X_train['teacher_number_of_previously_posted_projects'].values.reshape(1, -1))
prev_projects_cv = normalizer.transform(X_cv['teacher_number_of_previously_posted_projects'].values.reshape(1, -1))
prev_projects_test = normalizer.transform(X_test['teacher_number_of_previously_posted_projects'].values.reshape(1, -1))

print("After vectorizations")
print(prev_projects_train.shape, y_train.shape)
print(prev_projects_cv.shape, y_cv.shape)
print(prev_projects_test.shape, y_test.shape)
print("=="*100)
```

After vectorizations

(49041, 1) (49041,)

(24155, 1) (24155,)

(36052, 1) (36052,)

=====
=====

4) title word count

```
In [292]: normalizer = Normalizer()

normalizer.fit(X_train['title_word_count'].values.reshape(1, -1))

title_word_count_train = normalizer.transform(X_train['title_word_count'].values.reshape(1, -1))
title_word_count_cv = normalizer.transform(X_cv['title_word_count'].values.reshape(1, -1))
title_word_count_test = normalizer.transform(X_test['title_word_count'].values.reshape(1, -1))

print("After vectorizations")
print(title_word_count_train.shape, y_train.shape)
print(title_word_count_cv.shape, y_cv.shape)
print(title_word_count_test.shape, y_test.shape)
print("=="*100)
```

After vectorizations

(1, 49041) (49041,)

(1, 24155) (24155,)

(1, 36052) (36052,)

=====
=====

5) essay word count

```
In [293]: normalizer = Normalizer()

normalizer.fit(X_train['essay_word_count'].values.reshape(1,-1))

essay_word_count_train = normalizer.transform(X_train['essay_word_count'].values)
essay_word_count_cv = normalizer.transform(X_cv['essay_word_count'].values.reshape(1,-1))
essay_word_count_test = normalizer.transform(X_test['essay_word_count'].values.reshape(1,-1))

print("After vectorizations")
print(essay_word_count_train.shape, y_train.shape)
print(essay_word_count_cv.shape, y_cv.shape)
print(essay_word_count_test.shape, y_test.shape)
```

After vectorizations

```
(1, 49041) (49041,)
(1, 24155) (24155,)
(1, 36052) (36052,)
```

Assignment 4

1. Apply Multinomial NaiveBayes on these feature sets
 - : categorical, numerical features + project_title(BOW) + preprocessed_essay (BOW)
 - : categorical, numerical features + project_title(TFIDF)+ preprocessed_essay (TFIDF)
2. The hyper parameter tuning(find best Alpha)
 - : Find the best hyper parameter which will give the maximum AUC value
 - : Consider a wide range of alpha values for hyperparameter tuning, start as low as 0.00001
 - : Find the best hyper parameter using k-fold cross validation or simple cross validation data
 - : Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task of hyperparameter tuning
3. Feature importance
 - : Find the top 10 features of positive class and top 10 features of negative class for both feature sets and using values of feature_log_prob_ parameter of MultinomialNB and print their corresponding feature names
4. Representation of results
 - : You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure. Here on X-axis you will have alpha values, since they have a wide range, just to represent those alpha values on the graph, apply log function on those alpha values.
 - : Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data and plot the ROC curve on both train and test.
 - : Along with plotting ROC curve, you need to print the confusion matrix with predicted and original labels of test data points. Please visualize your confusion matrices using seaborn heatmaps.
5. Conclusion
 - : You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table please refer to this [prettytable library link](#)

2. Naive Bayes

Set 1: categorical, numerical features + project_title(BOW) + preprocessed_essay (BOW)

```
In [294]: # reshaping to appt shape so that it can be used with hstack
```

```
In [295]: price_train = (X_train['price'].values.reshape(-1,1))
price_cv = (X_cv['price'].values.reshape(-1,1))
price_test = (X_test['price'].values.reshape(-1,1))

quantity_train = (X_train['quantity'].values.reshape(-1,1))
quantity_cv = (X_cv['quantity'].values.reshape(-1,1))
quantity_test = (X_test['quantity'].values.reshape(-1,1))

prev_projects_train = (X_train['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
prev_projects_cv = (X_cv['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
prev_projects_test = (X_test['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))

title_word_count_train = (X_train['title_word_count'].values.reshape(-1,1))
title_word_count_cv = (X_cv['title_word_count'].values.reshape(-1,1))
title_word_count_test = (X_test['title_word_count'].values.reshape(-1,1))

essay_word_count_train = (X_train['essay_word_count'].values.reshape(-1,1))
essay_word_count_cv = (X_cv['essay_word_count'].values.reshape(-1,1))
essay_word_count_test = (X_test['essay_word_count'].values.reshape(-1,1))
```

```
In [296]: # merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039
from scipy.sparse import hstack

X_tr = hstack((categories_one_hot_train, sub_categories_one_hot_train, school_state_cat_train))
X_te = hstack((categories_one_hot_test, sub_categories_one_hot_test, school_state_cat_test))
X_cr = hstack((categories_one_hot_cv, sub_categories_one_hot_cv, school_state_cat_cv))
```

```
In [297]: print("Final Data matrix")
print(X_tr.shape, y_train.shape)
print(X_cr.shape, y_cv.shape)
print(X_te.shape, y_test.shape)
print("=="*100)
```

```
Final Data matrix
(49041, 7093) (49041,)
(24155, 7093) (24155,)
(36052, 7093) (36052,)
```

```
=====
=====
```

```
In [298]: def batch_predict(clf, data):
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability esti
# not the predicted outputs
y_data_pred = []
tr_loop = data.shape[0] - data.shape[0]%1000
# consider you X_tr shape is 49041, then your cr_loop will be 49041 - 49041%1000
# in this for loop we will iterate until the last 1000 multiplier
for i in range(0, tr_loop, 1000):
    y_data_pred.extend(clf.predict_proba(data[i:i+1000]))[:,1])
# we will be predicting for the last data points
y_data_pred.extend(clf.predict_proba(data[tr_loop:]))[:,1])
return y_data_pred
```

A) Random alpha values (hyperparameter)

```
In [299]: import matplotlib.pyplot as plt
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import roc_auc_score
import math

train_auc = []
cv_auc = []
log_alphas = []

alphas = [0.00001, 0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000]

for i in tqdm(alphas):
    nb = MultinomialNB(alpha = i, class_prior=[0.5,0.5])
    nb.fit(X_tr, y_train)
    y_train_pred = batch_predict(nb, X_tr)
    y_cv_pred = batch_predict(nb, X_cr)
    # roc_auc_score(y_true, y_score) the 2nd parameter should be probability esti
    # not the predicted outputs
    train_auc.append(roc_auc_score(y_train, y_train_pred))
    cv_auc.append(roc_auc_score(y_cv, y_cv_pred))

for a in tqdm(alphas):
    b = math.log(a)
    log_alphas.append(b)
```


```
100%|████████████████████████████████████████████████████████████████████████████████| 9/9 [00:01<00:00, 5.68it/s]
100%|████████████████████████████████████████████████████████████████████████████████| 9/9 [00:00<00:00, 9022.16it/s]
```



```
In [300]: plt.figure(figsize=(10,10))
plt.plot(log_alphas, train_auc, label='Train AUC')
plt.plot(log_alphas, cv_auc, label='CV AUC')

plt.scatter(log_alphas, train_auc, label='Train AUC points')
plt.scatter(log_alphas, cv_auc, label='CV AUC points')

plt.legend()
plt.xlabel("log alpha: hyperparameter")
plt.ylabel("AUC")
plt.title("alpha: hyperparameter v/s AUC")
plt.grid()
plt.show()
```



Summary

- As you can see in the plot, we've chosen hyperparameter alpha values as low as 0.0001 and as high as 1000. Since it's hard to plot the range provided, we've used log alphas on the x-axis and AUC on the y-axis.
- In addition to being able to display large ranges of values using log scales, log scales also prevent small values from being compressed into the graph's bottom.
- In both train and CV AUC lines, convergence can be observed at log alpha values closer to 7.
- We can see how both lines converge much more rapidly after alpha=10.

B) Gridsearch-cv using cv = 10 (K fold cross validation)

- return_train_score=True needs to be explicitly set for GridSearchCV function to return train scores
- Also verbose is set to 2 to display progress messages

```
In [301]: from sklearn.model_selection import GridSearchCV

nb = MultinomialNB(class_prior=[0.5,0.5])

parameters = {'alpha':[0.00001, 0.0001,0.001, 0.01, 0.1,0.5,0.8, 1, 10, 100, 1000]}

clf = GridSearchCV(nb, parameters, cv= 10, scoring='roc_auc',return_train_score=True)

clf.fit(X_tr, y_train)

train_auc= clf.cv_results_['mean_train_score']
train_auc_std= clf.cv_results_['std_train_score']
cv_auc = clf.cv_results_['mean_test_score']
cv_auc_std= clf.cv_results_['std_test_score']
```

Fitting 10 folds for each of 11 candidates, totalling 110 fits

```
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    6.8s
[Parallel(n_jobs=-1)]: Done 110 out of 110 | elapsed:    15.5s finished
```

```

In [302]: alphas = [0.00001, 0.0001, 0.001, 0.01, 0.1, 0.5, 0.8, 1, 10, 100, 1000]
log_alphas = []

for a in tqdm(alphas):
    b = math.log(a)
    log_alphas.append(b)

plt.figure(figsize=(20,10))

plt.plot(log_alphas, train_auc, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(log_alphas, train_auc - train_auc_std, train_auc + train_auc_std)

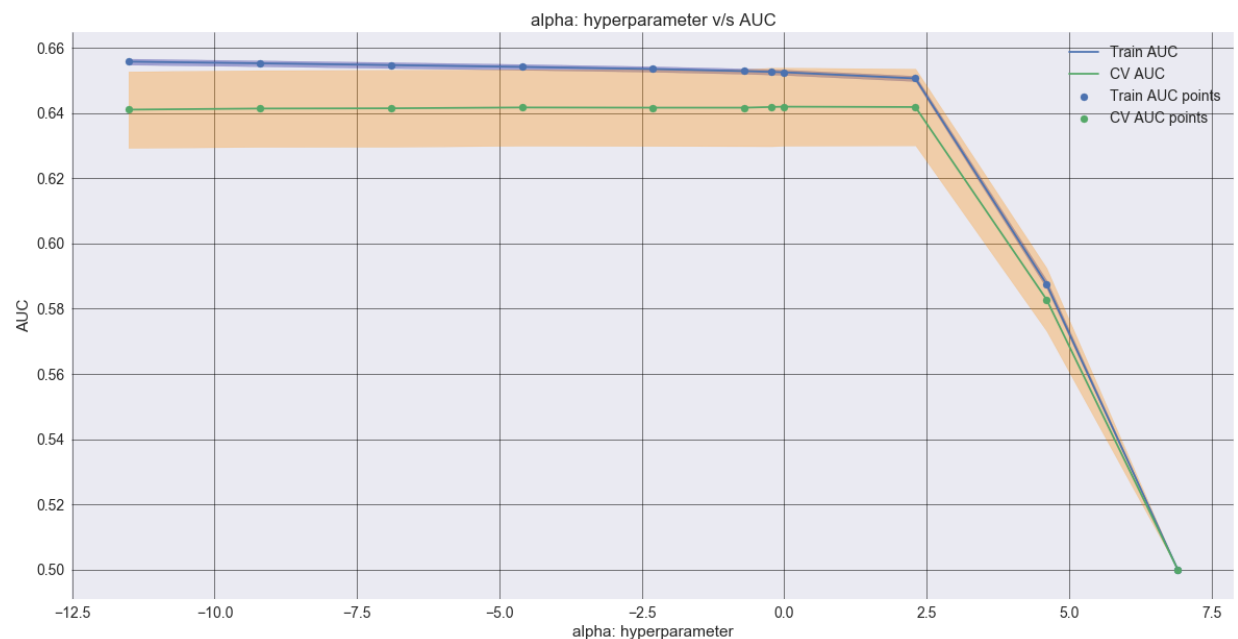
plt.plot(log_alphas, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(log_alphas, cv_auc - cv_auc_std, cv_auc + cv_auc_std, alpha=0.5)

plt.scatter(log_alphas, train_auc, label='Train AUC points')
plt.scatter(log_alphas, cv_auc, label='CV AUC points')

plt.legend()
plt.xlabel("alpha: hyperparameter")
plt.ylabel("AUC")
plt.title("alpha: hyperparameter v/s AUC")
plt.grid(color='black', linestyle='-', linewidth=0.5)
plt.show()

```

100% | 11/11 [00:00<?, ?it/s]



Summary

- We have started with hyperparameter alpha with as low as 0.0001 to 1000. Since it is difficult to plot the given range we have used log alphas on x-axis and AUC on y axis as shown in the plot.
- One of the main reason for using log scale is log scales allow a large range to be displayed without small values being compressed down into bottom of the graph.
- we observe that as log alpha approaches close to 7 ,both train AUC and cv AUC lines converge
- Using this plot we see after alpha=100 both lines converge at amuch higher rate

Train model using the best hyper-parameter value

- Using best_params_ attribute of gridsearch cv we can obtain the optimal value of alpha among the values we have selected
- It simplifes our task and we can be rest assured that selected hyperparameter is most optimal one

```
In [303]: print(clf.best_params_)
```

```
{'alpha': 1}
```

```
In [304]: best_alpha1=clf.best_params_
```

```
In [305]: # https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.htm
from sklearn.metrics import roc_curve, auc

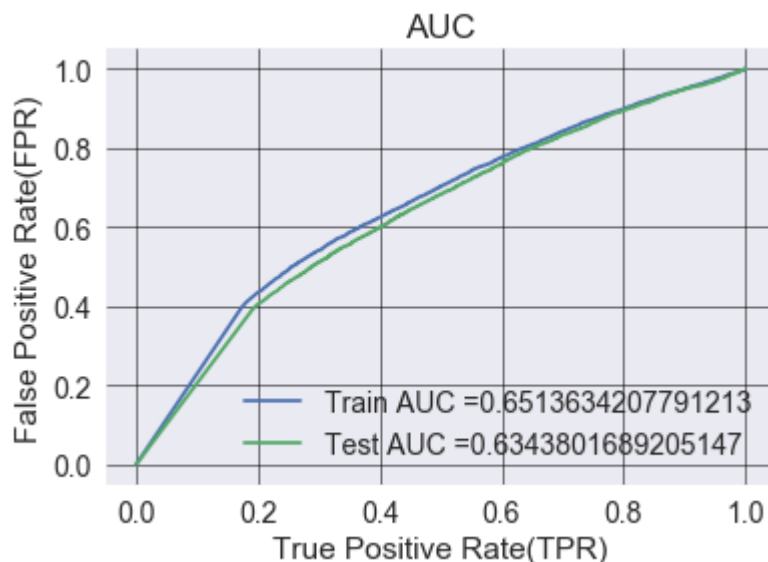
nb_bow = MultinomialNB(alpha = 1, class_prior=[0.5,0.5])

nb_bow.fit(X_tr, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimate
# not the predicted outputs

y_train_pred = batch_predict(nb_bow, X_tr)
y_test_pred = batch_predict(nb_bow, X_te)

train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)

plt.plot(train_fpr, train_tpr, label="Train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="Test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("True Positive Rate(TPR)")
plt.ylabel("False Positive Rate(FPR)")
plt.title("AUC")
plt.grid(color='black', linestyle='--', linewidth=0.5)
plt.show()
```



summary

- For Bow model for alpha= 1 ,we get train AUC of 0.65 and Test AUC of 0.63

D) confusion matrix

```
In [306]: def predict(proba, threshold, fpr, tpr):
            t = threshold[np.argmax(fpr*(1-tpr))]

            # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high

            print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold",
                  predictions = [])
            for i in proba:
                if i>=t:
                    predictions.append(1)
                else:
                    predictions.append(0)
            return predictions
```

train data

```
In [307]: print("="*100)
            from sklearn.metrics import confusion_matrix
            print("Train confusion matrix")
            print(confusion_matrix(y_train, predict(y_train_pred, tr_thresholds, train_fpr, t
```

```
=====
=====
Train confusion matrix
the maximum value of tpr*(1-fpr) 0.38243606289412657 for threshold 0.999
[[ 4137  3289]
 [14316 27299]]
```

```
In [308]: conf_matr_df_train_1 = pd.DataFrame(data=confusion_matrix(y_train, predict(y_train
```

```
the maximum value of tpr*(1-fpr) 0.38243606289412657 for threshold 0.999
```


```
In [309]: conf_matr_df_train_1
```

Out[309]:

	0	1
0	4137	3289
1	14316	27299

```
In [310]: sns.set(font_scale=1.4)#for label size  
sns.heatmap(conf_matr_df_train_1, annot=True,annot_kws={"size": 16}, fmt='.0f',li  
#mention prediction and actual
```

```
Out[310]: <matplotlib.axes._subplots.AxesSubplot at 0x1ba0e577f98>
```



Summary on train data

- In the following confusion matrix we observe that the model has 27268 true positives while true negatives are only 4174
- It has large number of false negatives which is 14347

test data

```
In [311]: print("=*100)
print("Test confusion matrix")
print(confusion_matrix(y_test, predict(y_test_pred, tr_thresholds, test_fpr, test_


=====
=====
Test confusion matrix
the maximum value of tpr*(1-fpr) 0.3629947374659989 for threshold 1.0
[[ 3497  1962]
 [13285 17308]]
```

```
In [312]: conf_matr_df_test_1 = pd.DataFrame(confusion_matrix(y_test, predict(y_test_pred,

the maximum value of tpr*(1-fpr) 0.3629947374659989 for threshold 1.0
```

```
In [313]: sns.set(font_scale=1.4)#for label size
sns.heatmap(conf_matr_df_test_1, annot=True,annot_kws={"size": 16}, fmt='g',linewidth
```

```
Out[313]: <matplotlib.axes._subplots.AxesSubplot at 0x1b99495fd68>
```



Summary on test data

- The number of true positives dominate ,there are 19080 in number,
- The least number among 4 quantites is false positives which are 2321 false positives
- similar trend is observed for false negatives which are roughly 11513

Set 2 : categorical, numerical features + project_title(TFIDF) + preprocessed_essay (TFIDF)


```
In [314]: # merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039
from scipy.sparse import hstack
X_tr = hstack((categories_one_hot_train, sub_categories_one_hot_train, school_state_cat
X_te = hstack((categories_one_hot_test, sub_categories_one_hot_test, school_state_cat
X_cr = hstack((categories_one_hot_cv, sub_categories_one_hot_cv, school_state_cat
```

```
In [315]: print("Final Data matrix")
print(X_tr.shape, y_train.shape)
print(X_cr.shape, y_cv.shape)
print(X_te.shape, y_test.shape)
print("="*100)
```

```
Final Data matrix
(49041, 7093) (49041,)
(24155, 7093) (24155,)
(36052, 7093) (36052,)
```

```
=====
=====
```

A) random alpha values

```
In [316]: train_auc = []
cv_auc = []
log_alphas = []

alphas = [0.00001, 0.0001, 0.001, 0.01, 0.1, 0.5, 0.8, 1, 10, 100, 1000]

for i in tqdm(alphas):
    nb = MultinomialNB(alpha = i, class_prior=[0.5, 0.5])
    nb.fit(X_tr, y_train)
    y_train_pred = batch_predict(nb, X_tr)
    y_cv_pred = batch_predict(nb, X_cr)
    # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates
    # not the predicted outputs
    train_auc.append(roc_auc_score(y_train, y_train_pred))
    cv_auc.append(roc_auc_score(y_cv, y_cv_pred))
for a in tqdm(alphas):
    b = math.log(a)
    log_alphas.append(b)
```


```
100%|████████████████████████████████████████████████████████████████████████████████|
11/11 [00:01<00:00, 5.60it/s]
100%|████████████████████████████████████████████████████████████████████████████████|
11/11 [00:00<00:00, 11019.19it/s]
```

```
In [317]: plt.figure(figsize=(20,10))

plt.plot(log_alphas, train_auc, label='Train AUC')
plt.plot(log_alphas, cv_auc, label='CV AUC')

plt.scatter(log_alphas, train_auc, label='Train AUC points')
plt.scatter(log_alphas, cv_auc, label='CV AUC points')

plt.legend()
plt.xlabel("alpha: hyperparameter")
plt.ylabel("AUC")
plt.title("alpha: hyperparameter v/s AUC")
plt.grid(color='black', linestyle='--', linewidth=0.5)
plt.show()
```



Summary

- Starting with an alpha between 0.0001 and 1000, it has been difficult to plot the given range on an x-axis. As depicted in the plot, log alphas are plotted on the x-axis and auc is plotted on the y-axis.
- In addition to allowing a large range of values to be displayed, log scales allow smaller values to be compressed into the bottom of the graph.
- The log alpha converges in both the train and cv AUC lines as log alpha approaches near 7.
- In this plot we can see that alpha=100 accelerates convergence of both lines at a much higher rate.

- The Cv Auc line remains constant for a long time in comparison to set1, only dropping after alpha=10.

```
In [318]: for it1,it2,it3 in zip(log_alphas, train_auc,cv_auc):
           print(2.71828**it1,it1,it2,it3)
```

```
1.0000077442305922e-05 -11.512925464970229 0.6305130724069298 0.619840639927975
4
0.00010000006195379677 -9.210340371976182 0.6300889526390047 0.6198141462169502
0.001000004646531159 -6.907755278982137 0.6296496306106348 0.6197826432468665
0.010000030976850406 -4.605170185988091 0.6292142660885494 0.6197096788196494
0.10000015488413208 -2.3025850929940455 0.6286031679012652 0.619643251108647
0.5000002331237219 -0.6931471805599453 0.6278726115939919 0.6197033355368737
0.8000001200785021 -0.2231435513142097 0.6275584040396465 0.619564997277391
1.0 0.0 0.6274713939116985 0.6195010308401471
9.999984511610785 2.302585092994046 0.6257835563122243 0.6184050503149723
99.99969023245559 4.605170185988092 0.5254896880103319 0.5252581182480224
999.9953534904313 6.907755278982137 0.5 0.5
```

B) Gridsearch-cv using cv = 10 (K fold cross validation)

```
In [319]: nb = MultinomialNB(class_prior=[0.5,0.5])

parameters = {'alpha':[0.00001, 0.0001, 0.001, 0.01, 0.1,0.25,0.5,0.8, 1,100]}

clf = GridSearchCV(nb, parameters, cv= 10, scoring='roc_auc',return_train_score=T

clf.fit(X_tr, y_train)

train_auc= clf.cv_results_['mean_train_score']
train_auc_std= clf.cv_results_['std_train_score']
cv_auc = clf.cv_results_['mean_test_score']
cv_auc_std= clf.cv_results_['std_test_score']
```

Fitting 10 folds for each of 10 candidates, totalling 100 fits

```
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    6.3s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:    13.1s finished
```

```
In [320]: alphas = [0.00001, 0.0001, 0.001, 0.01, 0.1, 0.25, 0.5, 0.8, 1, 100]
log_alphas = []

for a in tqdm(alphas):
    b = math.log(a)
    log_alphas.append(b)

plt.figure(figsize=(20,10))


plt.plot(log_alphas, train_auc, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(log_alphas, train_auc - train_auc_std, train_auc + train_auc_std, alpha=0.5)

plt.plot(log_alphas, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(log_alphas, cv_auc - cv_auc_std, cv_auc + cv_auc_std, alpha=0.5)

plt.scatter(log_alphas, train_auc, label='Train AUC points')
plt.scatter(log_alphas, cv_auc, label='CV AUC points')

plt.legend()
plt.xlabel("log alpha: hyperparameter")
plt.ylabel("AUC")
plt.title("alpha: hyperparameter v/s AUC")
plt.grid(color='black', linestyle='--', linewidth=0.5)
plt.show()
```

```
100%|██████████| 10/10 [00:00<00:00, 10022.23it/s]
```



```
In [321]: print( 2.67**0.6282283144644325)
```

```
1.8533030942243873
```

Summary

```
In [322]: for it1,it2,it3 in zip(log_alphas, train_auc,cv_auc):
           print(2.71828**it1,it1,it2)
```

```
1.0000077442305922e-05 -11.512925464970229 0.6314653388396043
0.0001000006195379677 -9.210340371976182 0.6309235800488118
0.001000004646531159 -6.907755278982137 0.6303640989349862
0.010000030976850406 -4.605170185988091 0.6297609865291701
0.10000015488413208 -2.3025850929940455 0.6290311619649992
0.25000023312377617 -1.3862943611198906 0.6286023244891439
0.5000002331237219 -0.6931471805599453 0.6282283144644325
0.8000001200785021 -0.2231435513142097 0.6278960124932909
1.0 0.0 0.6277296201884919
99.99969023245559 4.605170185988092 0.5049324818705895
```

- We have started with hyperparameter alpha with as low as 0.0001 to 1000. Since it is difficult to plot the given range we have used log alphas on x-axis and AUC on y axis as shown in the plot.
- One of the main reason for using log scale is log scales allow a large range to be displayed without small values being compressed down into bottom of the graph.
- we observe that as log alpha approaches close to 2 ,both train AUC and cv AUC lines converge
- Using this plot we see after alpha=100 both lines converge at a much higher rate
- One thing different from set1 plot is Cv AUC line remains constant for a long time ,only after log alpha=0.1 it starts dropping

C) Train model using the best hyper-parameter value of alpha

```
In [323]: # Code idea for best_params_ taken from here https://datascience.stackexchange.com
```

```
In [324]: best_alpha2=clf.best_params_
           print(best_alpha2)
```

```
{'alpha': 0.1}
```


```
In [359]: nb_tfidf = MultinomialNB(alpha = 0.1,class_prior=[0.5,0.5])

nb_tfidf.fit(X_tr, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimate
# not the predicted outputs

y_train_pred = batch_predict(nb_tfidf, X_tr)
y_test_pred = batch_predict(nb_tfidf, X_te)

train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)

plt.plot(train_fpr, train_tpr, label="Train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="Test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("True Positive Rate(TPR)")
plt.ylabel("False Positive Rate(FPR)")
plt.title("AUC")
plt.grid(color='black', linestyle='-', linewidth=0.5)
plt.show()
```



Summary

- From given plot we observe that we get train AUC of 0.652 and test AUC of 0.634

D) Confusion matrix

train data

```
In [360]: print("="*100)
print("Train confusion matrix")
print(confusion_matrix(y_train, predict(y_train_pred, tr_thresholds, train_fpr, t


=====
=====
Train confusion matrix
the maximum value of tpr*(1-fpr) 0.3830785282826924 for threshold 0.999
[[ 4129  3297]
 [14204 27411]]
```

```
In [361]: conf_matr_df_train_2 = pd.DataFrame(confusion_matrix(y_train, predict(y_train_pre

the maximum value of tpr*(1-fpr) 0.3830785282826924 for threshold 0.999
```

```
In [362]: sns.set(font_scale=1.4)#for label size
sns.heatmap(conf_matr_df_train_2, annot=True,annot_kws={"size": 16}, fmt='g',line
```

```
Out[362]: <matplotlib.axes._subplots.AxesSubplot at 0x1b990ccbda0>
```



Summary

- For training data we get roughly 27411 true positives and 4129 true negatives and false positives each
- Again we have roughly 14204 false negatives which are alot in number

test data

```
In [363]: print("=*100)
print("Test confusion matrix")
print(confusion_matrix(y_test, predict(y_test_pred, tr_thresholds, test_fpr, test_


=====
=====
Test confusion matrix
the maximum value of tpr*(1-fpr) 0.363407001160974 for threshold 1.0
[[ 3480 1979]
 [13193 17400]]
```

```
In [364]: conf_matr_df_test_2 = pd.DataFrame(confusion_matrix(y_test, predict(y_test_pred,

the maximum value of tpr*(1-fpr) 0.363407001160974 for threshold 1.0
```

```
In [365]: sns.set(font_scale=1.4)#for label size
sns.heatmap(conf_matr_df_test_2, annot=True,annot_kws={"size": 16}, fmt='g',linewidth
```

```
Out[365]: <matplotlib.axes._subplots.AxesSubplot at 0x1b990190780>
```



summary

- we have roughly 17400 true positives for test data and roughly 3480 true negatives
- Again false negatives are pretty high in number(13k)

Select best 30 features of both Positive and negative class for both the sets of data

set1 Bow


```
In [392]: X_tr = hstack((categories_one_hot_train, sub_categories_one_hot_train, school_state_cat
X_te = hstack((categories_one_hot_test, sub_categories_one_hot_test, school_state_cat
X_cr = hstack((categories_one_hot_cv, sub_categories_one_hot_cv, school_state_cat
```

```
In [443]: nb_bow = MultinomialNB(alpha = 1,class_prior=[0.5,0.5])

nb_bow.fit(X_tr, y_train)
```

```
Out[443]: MultinomialNB(alpha=1, class_prior=[0.5, 0.5], fit_prior=True)
```

```
In [444]: # Collecting feature names for BOW set1
# adding to end of list by concatenating features
# Code snippet taken from here https://stackabuse.com/append-vs-extend-in-python-

bow_features_names1 = []
```

```
In [445]: for cnt in vectorizer_proj.get_feature_names() :
            bow_features_names1.append(cnt)

for cnt1 in vectorizer_sub_proj.get_feature_names() :
            bow_features_names1.append(cnt1)

for cnt2 in vectorizer_states.get_feature_names() :
            bow_features_names1.append(cnt2)

for cnt3 in vectorizer_grade.get_feature_names() :
            bow_features_names1.append(cnt3)

for cnt4 in vectorizer_teacher.get_feature_names() :
            bow_features_names1.append(cnt4)
```

```
In [446]: len(bow_features_names1)
```

```
Out[446]: 98
```

```
In [447]: bow_features_names1.append("price")

bow_features_names1.append("quantity")

bow_features_names1.append("prev_proposed_projects")

bow_features_names1.append("title_word_count")

bow_features_names1.append("essay_word_count")

len(bow_features_names1)
```

```
Out[447]: 103
```

```
In [448]: for cnt5 in vectorizer_bow_title.get_feature_names() :  
          bow_features_names1.append(cnt5)
```

```
In [449]: len(bow_features_names1)
```

```
Out[449]: 2093
```

```
In [450]: for cnt6 in vectorizer_bow_essay.get_feature_names() :  
          bow_features_names1.append(cnt6)
```

```
In [451]: len(bow_features_names1)
```

```
Out[451]: 7093
```

Top 30 positive features BOW

NOTE

- Using argsort by default it sorts in ascending order, but we need sorted log probabilities in descending order.
- While sorting the log probabilities in ascending order we get least imp features because just for an example say f1 has prob -16 and f2 has prob -14 .The actual prob of f1 would be $\exp(-16)$ and actual prob of f2 would be $\exp(-14)$. Clearly f2 has higher prob so it is an imp feature

```
In [452]: # To use argsort for descending order  
          # Code snippet taken from https://stackoverflow.com/questions/16486252/is-it-poss
```

```
In [453]: pos_class_prob_sorted = nb_bow.feature_log_prob_[1, :].argsort()[::-1]
for i in pos_class_prob_sorted[:30]:
    print(bow_features_names1[i])
```

```
price
essay_word_count
quantity
prev_proposed_projects
students
title_word_count
school
learning
classroom
not
learn
help
grades
many
nannan
need
reading
work
use
love
able
day
come
class
would
technology
also
books
skills
year
```

30 negative features from BOW set1

```
In [454]: # To use argsort for descending order
# Code snippet taken from https://stackoverflow.com/questions/16486252/is-it-poss
# using log_prob_ Code taken from https://scikit-learn.org/stable/modules/generato
```

```
In [455]: neg_class_prob_sorted = nb_bow.feature_log_prob_[0,:].argsort()[::-1]
          for i in neg_class_prob_sorted[0:30]:
              print(bow_features_names1[i])
```

```
price
essay_word_count
quantity
prev_proposed_projects
students
title_word_count
school
learning
classroom
not
learn
help
grades
nannan
many
need
work
come
love
reading
materials
able
skills
class
day
use
want
year
make
new
```

Summary

- Words like learn is present in negative class but not in positive class
- Few words are similar but their relative ordering is different between the two sets

tfidf top features

```
In [456]: X_tr1 = hstack((categories_one_hot_train, sub_categories_one_hot_train, school_state_cat
X_te1 = hstack((categories_one_hot_test, sub_categories_one_hot_test, school_state_cat
X_cr1 = hstack((categories_one_hot_cv, sub_categories_one_hot_cv, school_state_cat
```

```
In [432]: X_tr1.shape
```

```
Out[432]: (49041, 7093)
```

```
In [457]: nb_tfidf = MultinomialNB(alpha = 0.1, class_prior=[0.5, 0.5])  
nb_tfidf.fit(X_tr, y_train)
```

```
Out[457]: MultinomialNB(alpha=0.1, class_prior=[0.5, 0.5], fit_prior=True)
```

```
In [458]: tfidf_features_names = []
```

```
In [459]: for ct1 in vectorizer_proj.get_feature_names() :  
          tfidf_features_names.append(ct1)  
  
          for ct2 in vectorizer_sub_proj.get_feature_names() :  
            tfidf_features_names.append(ct2)  
  
          for ct3 in vectorizer_states.get_feature_names() :  
            tfidf_features_names.append(ct3)  
  
          for ct4 in vectorizer_grade.get_feature_names() :  
            tfidf_features_names.append(ct4)  
  
          for ct5 in vectorizer_teacher.get_feature_names() :  
            tfidf_features_names.append(ct5)
```

```
In [460]: len(tfidf_features_names)
```

```
Out[460]: 98
```

```
In [461]: tfidf_features_names.append("price")  
          tfidf_features_names.append("quantity")  
          tfidf_features_names.append("prev_proposed_projects")  
          tfidf_features_names.append("title_word_count")  
          tfidf_features_names.append("essay_word_count")  
  
          for ct6 in vectorizer_tfidf_titles.get_feature_names() :  
            tfidf_features_names.append(ct6)  
  
          for ct7 in vectorizer_tfidf_essay.get_feature_names() :  
            tfidf_features_names.append(ct7)
```

```
In [462]: len(tfidf_features_names)
```

```
Out[462]: 7093
```

positive features of tfidf

```
In [463]: pos_class_prob_sorted_tfidf = nb_tfidf.feature_log_prob_[1, :].argsort()[::-1][:7]
for i in pos_class_prob_sorted_tfidf[0:30]:
    print(tfidf_features_names[i])
```

```
price
essay_word_count
quantity
prev_proposed_projects
students
title_word_count
school
learning
classroom
not
learn
help
grades
many
nannan
need
reading
work
use
love
able
day
come
class
would
technology
also
books
skills
year
```

Negative features from Tfidf

```
In [464]: neg_class_prob_sorted_tfidf = nb_tfidf.feature_log_prob_[0, :].argsort()[::-1][:7]
for i in neg_class_prob_sorted_tfidf[0:30]:
    print(tfidf_features_names[i])
```

```
price
essay_word_count
quantity
prev_proposed_projects
students
title_word_count
school
learning
classroom
not
learn
help
grades
nannan
many
need
work
come
love
reading
materials
able
skills
class
day
use
want
year
make
new
```

summary

- Again important features appear similar at first glance but actually there are some differences compared to set1
- relative ordering is different between the two sets

conclusions

```
In [465]: # Please compare all your models using Prettytable Library
# http://zetcode.com/python/prettytable/

from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Vectorizer", "Model", "Alpha:Hyper Parameter", "Test AUC"]

x.add_row(["BOW", "Naive Bayes", 1, 0.634380])
x.add_row(["TFIDF", "Naive Bayes", 0.1, 0.634361])

print(x)
```

Vectorizer	Model	Alpha:Hyper Parameter	Test AUC
BOW	Naive Bayes	1	0.63438
TFIDF	Naive Bayes	0.1	0.634361

- We conclude that Naive bayes gives better AUC than KNN
- also it is very fast as compared to KNN.
- naive Bayes is super interpretable because of probability values, we can get feature importance very easily as seen above
- There is strong possibility that Naive bayes can overfit if alpha has not been found properly